

## Instalación de Hadoop single node pseudo-distribuida

Partimos de que se tiene instalado Hadoop single node. Para la última parte también se necesita el programa que calcula la temperatura máxima

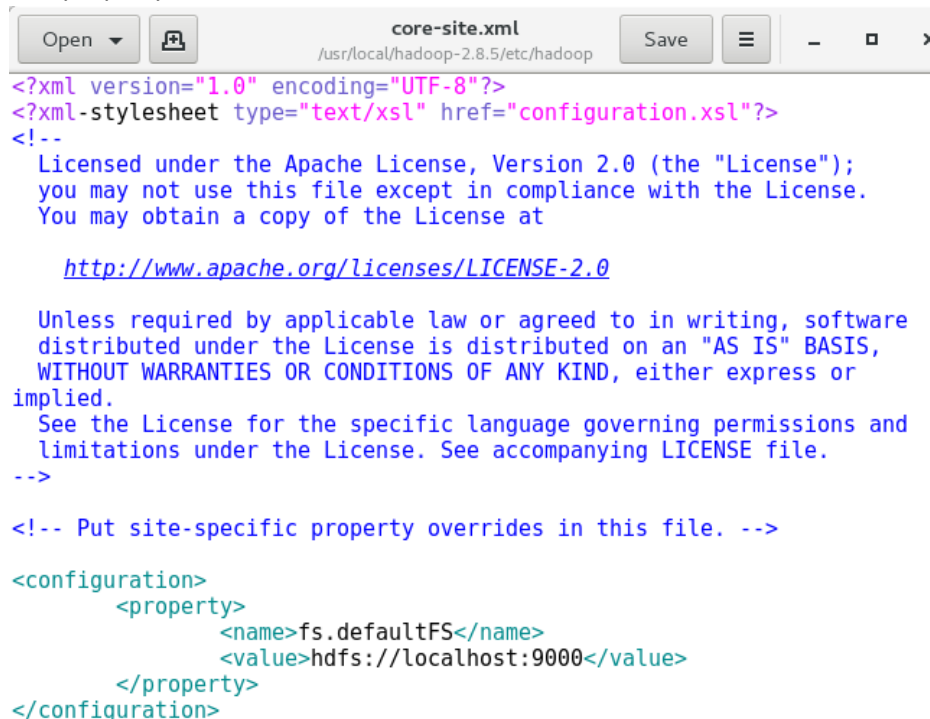
### 1. Modificamos la configuración de Hadoop:

#### 1.1. Ejecutamos: gedit \$HADOOP\_HOME/etc/hadoop/core-site.xml

```
[moranjesus@localhost ~]$ gedit $HADOOP_HOME/etc/hadoop/core-site.xml
```

#### 1.2. Añadimos dentro de configuración:

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
```

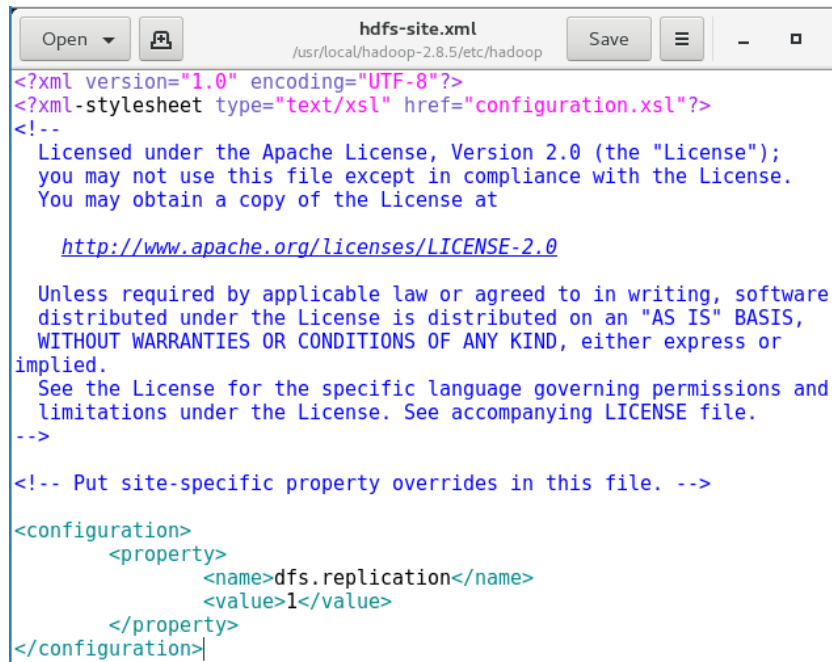


#### 1.3. Ejecutamos: gedit \$HADOOP\_HOME/etc/hadoop/hdfs-site.xml

```
[moranjesus@localhost ~]$ gedit $HADOOP_HOME/etc/hadoop/hdfs-site.xml
```

#### 1.4. Añadimos dentro de configuración:

```
<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>
```



```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
  implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
```

- 1.5. Ejecutamos: `cp $HADOOP_HOME/etc/hadoop/mapred-site.xml.template $HADOOP_HOME/etc/hadoop/mapred-site.xml`

Nota: puede que ya esté copiado y no se necesite copiar

```
[moranjesus@localhost ~]$ cp $HADOOP_HOME/etc/hadoop/mapred-site.xml.template $HADOOP_HOME/etc/hadoop/
mapred-site.xml
```

- 1.6. Ejecutamos: `gedit $HADOOP_HOME/etc/hadoop/mapred-site.xml`

```
[moranjesus@localhost ~]$ gedit $HADOOP_HOME/etc/hadoop/mapred-site.xml
```

- 1.7. Añadimos dentro de configuración:

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
<property>
  <name>yarn.app.mapreduce.am.env</name>
  <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
  <name>mapreduce.map.env</name>
  <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
  <name>mapreduce.reduce.env</name>
  <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
  <name>mapreduce.application.classpath</name>
  <value>$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/*:$HADOOP_MAPRE
D_HOME/share/hadoop/mapreduce/lib/*</value>
</property>
```



```
http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

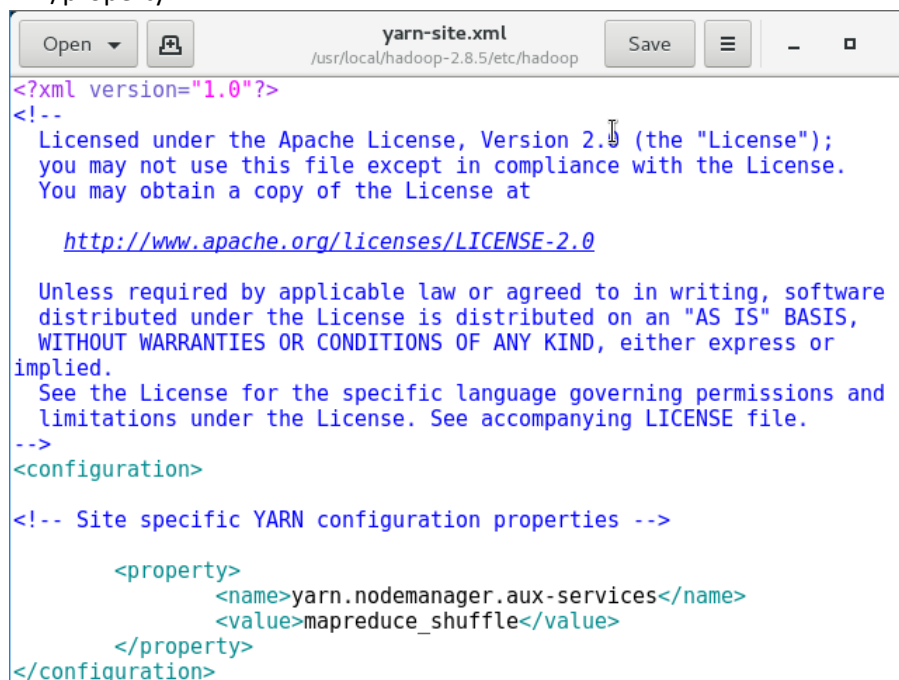
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>yarn.app.mapreduce.am.env</name>
    <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
  </property>
  <property>
    <name>mapreduce.map.env</name>
    <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
  </property>
  <property>
    <name>mapreduce.reduce.env</name>
    <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
  </property>
  <property>
    <name>mapreduce.application.classpath</name>
    <value>$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/*:$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/
lib/*</value>
  </property>
</configuration>
```

1.8. Ejecutamos: `gedit $HADOOP_HOME/etc/hadoop/yarn-site.xml`

`[moranjesus@localhost ~]$ gedit $HADOOP_HOME/etc/hadoop/yarn-site.xml`

1.9. Añadimos dentro de configuración:

```
</property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
```



```
<?xml version="1.0"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->
<configuration>

<!-- Site specific YARN configuration properties -->

  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

2. Formateamos el sistema de archivos: `hdfs namenode -format`

```
[moranjesus@localhost ~]$ hdfs namenode -format
```

3. Hadoop tiene un “problema” con las últimas versiones de java:  
<https://issues.apache.org/jira/browse/HADOOP-10848>. Hay varias formas de ejecutar Hadoop, lo recomendado es utilizar java 8. Para ello:

3.1. Nos conectamos como sudo: su -

3.2. Instalamos el openjdk 8, para ello: dnf install java-1.8.0-openjdk-devel.x86\_64

```
[moranjesus@localhost ~]$ su -  
Password:  
[root@localhost ~]# dnf install java-1.8.0-openjdk-devel.x86_64
```

3.3. Seleccionamos el nuevo java, para ello: sudo alternatives --config java

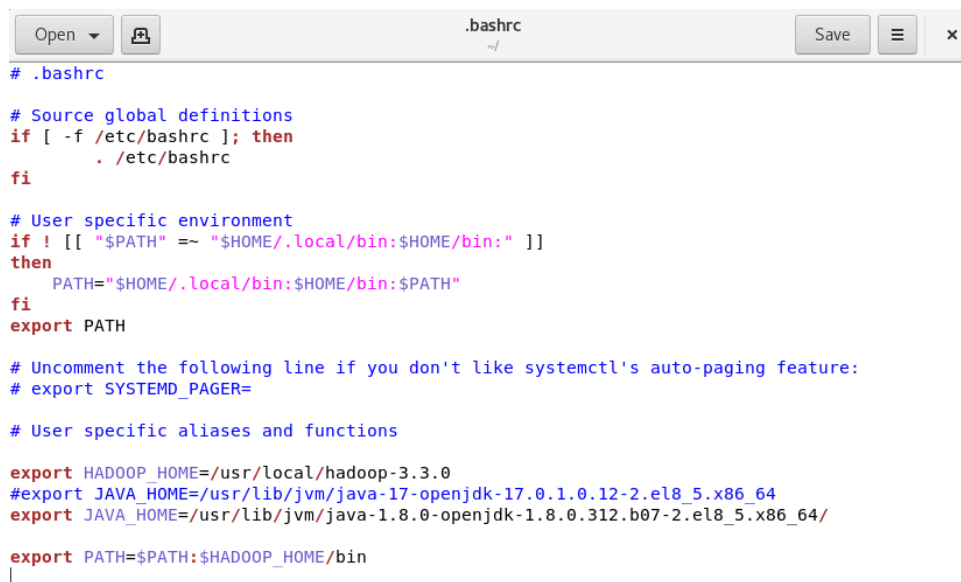
```
[root@localhost ~]# sudo alternatives --config java  
There are 2 programs which provide 'java'.  
  
Selection    Command  
-----  
+ 1          java-17-openjdk.x86_64 (/usr/lib/jvm/java-17-openjdk-17.0.1.0.12-2.el8_5.x86_64/bin/java)  
* 2          java-1.8.0-openjdk.x86_64 (/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.312.b07-2.el8_5.x86_64/jre/bin/java)  
  
Enter to keep the current selection[+], or type selection number: 2  
[root@localhost ~]#
```

Importante seleccionar el java 8 que acabamos de descargar

3.4. Salimos como usuario root, para ello: exit

```
[root@localhost ~]# exit  
logout  
[moranjesus@localhost ~]$
```

3.5. Como nuestro usuario editamos la variable de entorno de JAVA\_HOME para poner la nueva ubicación de java. Para ello: gedit ~/.bashrc



```
Open .bashrc Save  
# .bashrc  
  
# Source global definitions  
if [ -f /etc/bashrc ]; then  
    . /etc/bashrc  
fi  
  
# User specific environment  
if ! [ "$PATH" = "$HOME/.local/bin:$HOME/bin:" ]  
then  
    PATH="$HOME/.local/bin:$HOME/bin:$PATH"  
fi  
export PATH  
  
# Uncomment the following line if you don't like systemctl's auto-paging feature:  
# export SYSTEMD_PAGER=  
  
# User specific aliases and functions  
  
export HADOOP_HOME=/usr/local/hadoop-3.3.0  
#export JAVA_HOME=/usr/lib/jvm/java-17-openjdk-17.0.1.0.12-2.el8_5.x86_64  
export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.312.b07-2.el8_5.x86_64/  
  
export PATH=$PATH:$HADOOP_HOME/bin
```

Notar que hemos comentado la anterior JAVA\_HOME y puesto la nueva. Puede que los números del final sean diferentes para cada persona. Para comprobar cuáles son los números podéis hacer ls /usr/lib/jvm y copiar los números adecuados

3.6. Ahora hacemos un source para tener ya disponible la nueva variable de entorno, para ello: source ~/.bashrc

```
[moranjesus@localhost ~]$ source ~/.bashrc
```

#### 4. Inicializamos los servicios:

##### 4.1. Inicializamos el sistema de archivos distribuido: \$HADOOP\_HOME/sbin/start-dfs.sh

```
[moranjesus@localhost ~]$ $HADOOP_HOME/sbin/start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [localhost.localdomain]
localhost.localdomain: Warning: Permanently added 'localhost.localdomain' (ECDSA
) to the list of known hosts.
[moranjesus@localhost ~]$
```

##### 4.2. Inicializamos yarn: \$HADOOP\_HOME/sbin/start-yarn.sh

```
[moranjesus@localhost ~]$ $HADOOP_HOME/sbin/start-yarn.sh
Starting resourcemanager
Starting nodemanagers
[moranjesus@localhost ~]$
```

##### 4.3. Entramos en la página <http://localhost:8088> (del equipo almaLinux)

Applications Places Firefox Web Browser es Thu 02:43

All Applications - Mozilla Firefox

All Applications

localhost:8088/cluster

**hadoop** All Applications

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total
0	0	0	0	0	0 B	8 GB

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes
1	0	0	0

Scheduler Metrics

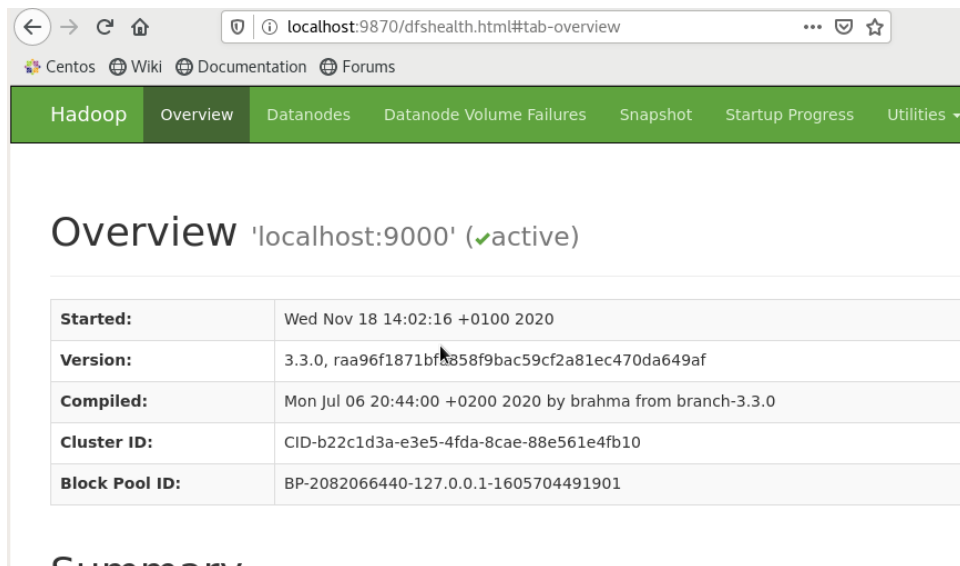
Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:1024, vCores:1>

Show 20 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Container
No data available in table										

Showing 0 to 0 of 0 entries

##### 4.4. Entramos en la página <http://localhost:9870> (del equipo almalinux)



- 4.5. Comprobamos que están los daemons ejecutándose: jps

```
[moranjesus@localhost ~]$ jps
16356 NodeManager
16884 Jps
15589 NameNode
16245 ResourceManager
15944 SecondaryNameNode
15727 DataNode
[moranjesus@localhost ~]$
```

5. Creamos una estructura de carpetas en HDFS:

- 5.1. Creamos la carpeta user, para ello: `hdfs dfs -mkdir /user`

```
[moranjesus@localhost ~]$ hdfs dfs -mkdir /user
[moranjesus@localhost ~]$
```

- 5.2. Creamos la carpeta de nuestro usuario: `hdfs dfs -mkdir /user/moranjesus`

```
[moranjesus@localhost ~]$ hdfs dfs -mkdir /user/moranjesus
[moranjesus@localhost ~]$
```

6. Subimos datos a HDFS:

- 6.1. Ejecutamos la subida:


```
hdfs dfs -put /home/moranjesus/Desktop/tempMax/medidas.txt
misDatosEnHDFS
[moranjesus@localhost ~]$ hdfs dfs -put /home/moranjesus/Desktop/tempMax/medidas.txt misDatosEnHDFS
[moranjesus@localhost ~]$
```

- 6.2. Comprobamos que los datos se introdujeron en HDFS, para ello:





```
hdfs dfs -ls /user/moranjesus
nota: equivale a hdfs dfs -ls
[moranjesus@localhost ~]$ hdfs dfs -ls /user/moranjesus
Found 1 items
-rw-r--r-- 1 moranjesus supergroup 151 2018-11-22 03:02 /user/moranjesus/misDatosEnHDFS
[moranjesus@localhost ~]$
```

- 6.3. Comprobamos desde el servidor web HDFS que se introdujeron los datos, para ello:



Desde el navegador web: <http://localhost:9870> (de la máquina almaLinux)  
Luego, vamos a Utilities -> browse the file system:



## Browse Directory


/ Go!    

Show 25 entries Search:





	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	drwxr-xr-x	<a href="#">moranjesus</a>	<a href="#">supergroup</a>	0 B	Nov 18 14:13	0	0 B	<a href="#">user</a> 

Showing 1 to 1 of 1 entries Previous 1 Next



## Entramos en /user/moranjesus



## Browse Directory

/user/moranjesus Go!    

Show 25 entries Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	<a href="#">moranjesus</a>	<a href="#">supergroup</a>	151 B	Nov 18 14:13	1	128 MB	<a href="#">misDatosEnHDFS</a> 

Showing 1 to 1 of 1 entries Previous 1 Next

### 7. Ejecutamos un programa sobre los datos de HDFS:

#### 7.1. Nos ubicamos en la carpeta que tenemos el programa maxTemp, para ello: cd

/home/moranjesus/Desktop/tempMax

```
[moranjesus@localhost tempMax]$
```

#### 7.2. Ejecutamos el programa:

```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.3.0.jar -files ./mapperMaxTemp.py,./reducerMaxTemp.py -mapper ./mapperMaxTemp.py -reducer ./reducerMaxTemp.py -combiner ./reducerMaxTemp.py -input misDatosEnHDFS -output salidaEnHDFS
```

```
[moranjesus@localhost tempMax]$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.3.0.jar -files ./mapperMaxTemp.py,./reducerMaxTemp.py -mapper ./mapperMaxTemp.py -reducer ./reducerMaxTemp.py -combiner ./reducerMaxTemp.py -input misDatosEnHDFS -output salidaEnHDFS
```

```

File System Counters
  FILE: Number of bytes read=39
  FILE: Number of bytes written=485885
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=435
  HDFS: Number of bytes written=27
  HDFS: Number of read operations=9
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Killed map tasks=1
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=22310
  Total time spent by all reduces in occupied slots (ms)=5428
  Total time spent by all map tasks (ms)=22310
  Total time spent by all reduce tasks (ms)=5428
  Total vcore-milliseconds taken by all map tasks=22310
  Total vcore-milliseconds taken by all reduce tasks=5428
  Total megabyte-milliseconds taken by all map tasks=22845440
  Total megabyte-milliseconds taken by all reduce tasks=5558272
Map-Reduce Framework
  Map input records=11
  Map output records=11
  Map output bytes=77
  Map output materialized bytes=45
  Input split bytes=208
  Combine input records=11
  Combine output records=3
  Reduce input groups=3

```

7.3. Comprobamos desde la interfaz gráfica que se ejecutó el programa:

<http://localhost:8088>

The screenshot displays the Hadoop cluster management interface. The sidebar on the left contains a 'Cluster' section with links for 'About', 'Nodes', 'Node Labels', 'Applications', 'Scheduler', and 'Tools'. The main content area is divided into three sections: 'Cluster Metrics', 'Cluster Nodes Metrics', and 'Scheduler Metrics'. The 'Cluster Metrics' section shows a table with columns for Apps Submitted, Apps Pending, Apps Running, Apps Completed, and Containers. The 'Cluster Nodes Metrics' section shows a table with columns for Active Nodes, Decommissioning Nodes, and Decommissioned. The 'Scheduler Metrics' section shows a table with columns for Scheduler Type, Scheduling Resource Type, and Minim. The 'Scheduler Metrics' table shows a single application running with ID 'application\_1542850911288\_0001', User 'moranjesus', Name 'streamjob3389107036576450975.jar', Application Type 'MAPREDUCE', and Queue 'default'.

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers
1	0	1	0	0

Active Nodes	Decommissioning Nodes	Decommissioned
1	0	0

Scheduler Type	Scheduling Resource Type	Minim
Capacity Scheduler	[MEMORY]	<memory:1024, vCor

ID	User	Name	Application Type	Queue
application_1542850911288_0001	moranjesus	streamjob3389107036576450975.jar	MAPREDUCE	default

7.4. Si entramos dentro de la application\_\* podemos ver más información:



Application application\_1... x +

localhost:8088/cluster/app/application\_1542850911288\_0001

Logged in as: dr.who

## Application application\_1542850911288\_0001

Cluster

- About
- Nodes
- Node Labels
- Applications
  - NEW
  - NEW SAVING
  - SUBMITTED
  - ACCEPTED
  - RUNNING
  - FINISHED
  - FAILED
  - KILLED
- Scheduler

Tools

Kill Application

Application Overview

User:	<a href="#">moranjesus</a>
Name:	streamjob3389107036576450975.jar
Application Type:	MAPREDUCE
Application Tags:	
Application Priority:	0 (Higher Integer value indicates higher priority)
YarnApplicationState:	FINISHED
Queue:	default
FinalStatus Reported by AM:	SUCCEEDED
Started:	Thu Nov 22 03:21:26 +0100 2018
Elapsed:	33sec
Tracking URL:	<a href="#">History</a>
Log Aggregation Status:	DISABLED
Diagnostics:	
Unmanaged Application:	false
Application Node Label expression:	<Not set>
AM container Node Label expression:	<DEFAULT_PARTITION>

Application Metrics

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0

7.5. Comprobamos que la salida se creó en HDFS: `hdfs dfs -ls`

```
[moranjesus@localhost tempMax]$ hdfs dfs -ls
Found 2 items
-rw-r--r-- 1 moranjesus supergroup 150 2021-11-20 14:31 misDatosEnHDFS
drwxr-xr-x - moranjesus supergroup 0 2021-11-20 14:37 salidaEnHDFS
[moranjesus@localhost tempMax]$
```

7.6. Comprobamos desde la interfaz gráfica que se creó: `http://localhost:9870` luego ir a utilities -> browse the filesystem -> /user/moranjesus

localhost:9870/explorer.html#/user/moranjesus

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

## Browse Directory

/user/moranjesus Go!

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	<a href="#">moranjesus</a>	<a href="#">supergroup</a>	151 B	Nov 18 14:13	1	128 MB	<a href="#">misDatosEnHDFS</a>
drwxr-xr-x	<a href="#">moranjesus</a>	<a href="#">supergroup</a>	0 B	Nov 18 16:49	0	0 B	<a href="#">salidaEnHDFS</a>

Showing 1 to 2 of 2 entries Previous 1 Next

7.7. Si queremos ver los datos de salida podemos descargarlos desde la interfaz gráfica, o también mostrarlos por la terminal (en este caso podemos porque son pocos, si fuesen GBs de salida, no podríamos): `hdfs dfs -cat salidaEnHDFS/*`

```
[moranjesus@localhost tempMax]$ hdfs dfs -cat salidaEnHDFS/*
1999      5.0
2000      6.0
2001      3.0
[moranjesus@localhost tempMax]$ █
```

8. Cuando instalamos hadoop distribuido (o pseudodistribuido), el sistema de archivos por defecto deja de ser el local y pasa a ser hdfs. Además, todas las consultas irán relativas al usuario, es decir, si busco miEntrada, realmente estará buscando hdfs:/user/moranjesus/miEntrada (si es un archivo, el archivo, y si es una carpeta, los archivos de esa carpeta). Pero en ocasiones puede interesarnos ejecutar algo con datos del sistema de archivos local. Para ello podemos ejecutarlos utilizando el acrónimo file:/ Por ejemplo:

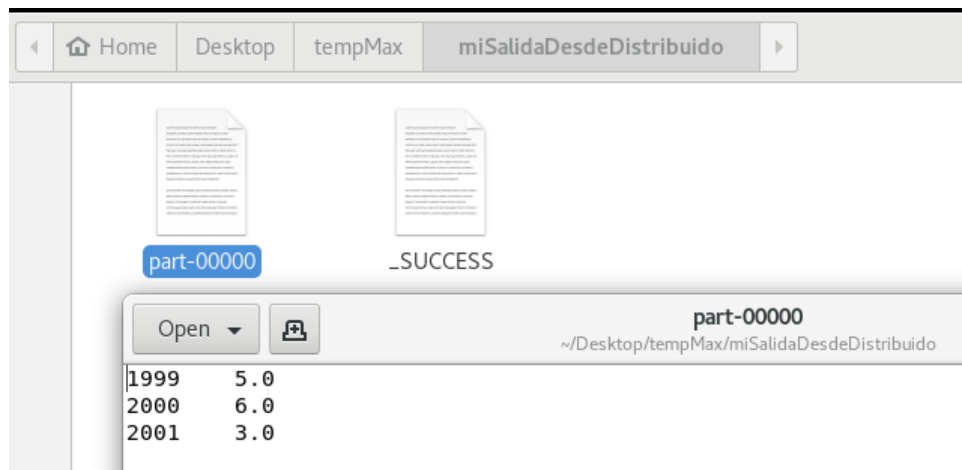
```
hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.3.0.jar -
files ./mapperMaxTemp.py,./reducerMaxTemp.py -mapper ./mapperMaxTemp.py -
reducer ./reducerMaxTemp.py -combiner ./reducerMaxTemp.py -input
file:/home/moranjesus/Desktop/tempMax/medidas.txt -output
file:/home/moranjesus/Desktop/tempMax/miSalidaDesdeDistribuido
```

```
[moranjesus@localhost tempMax]$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/h
adoop-streaming-3.3.0.jar -files ./mapperMaxTemp.py,./reducerMaxTemp.py -mapper
./mapperMaxTemp.py -reducer ./reducerMaxTemp.py -combiner ./reducerMaxTemp.py -i
nput file:./medidas.txt -output file:./miSalidaDesdeDistribuido █
```

- 8.1. Tras ejecutar, tendremos:

```
File System Counters
  FILE: Number of bytes read=264
  FILE: Number of bytes written=805778
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=202
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=2
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
```

- 8.2. Los datos se guardaron en nuestro sistema de archivos local (en una carpeta de nuestro ordenador):



9. Si queremos apagar los servicios de Hadoop: `$HADOOP_HOME/sbin/stop-all.sh`  
(también se pueden apagar de uno en uno)