

Procesamiento de datos masivos

Jesús Morán

- R vs SparkR
- Instalación SparkR
- Operaciones SparkR
- Visualización de datos
- Machine Learning

■ R:


- ☐ No soporta grandes cantidades de datos
- ☐ Muy utilizado por analistas de datos
- ☐ Muchas librerías
- ☐ `data.frame`

■ SparkR:

- Soporta grandes cantidades de datos
- Tiene objeto DataFrame
 - ~ data.frame de R pero limitado
 - Varias Row + Schema
 - Es distribuido
 - Se puede convertir en data.frame
- Hay operaciones Spark que no se pueden hacer con SparkR

- Instalar Spark
- Instalar R: (en todos los servidores del cluster)
 - Repositorio de R: `yum install epel-release`
 - Instalar R: `yum install R`
- Instalar librerías: (en todos los servidores del cluster)
 - Entrar en R e instalar
 - `> install.packages("ggplot2")`
 - `> install.packages("knitr")`
 - `> install.packages("devtools")`

■ SparkR

 hadmin@MBDUIV0:~

```
[hadmin@MBDUIV0 ~]$ sparkR
```

```
R version 3.5.1 (2018-07-02) -- "Feather Spray"
Copyright (C) 2018 The R Foundation for Statistical Computing
Platform: x86_64-redhat-linux-gnu (64-bit)
```

R es un software libre y viene sin GARANTIA ALGUNA.
Usted puede redistribuirlo bajo ciertas circunstancias.
Escriba 'license()' o 'licence()' para detalles de distribucion.

R es un proyecto colaborativo con muchos contribuyentes. Escriba 'contributors()' para obtener más información y 'citation()' para saber cómo citar R o paquetes de R en publicaciones.

Escriba 'demo()' para demostraciones, 'help()' para el sistema on-line de ayuda, o 'help.start()' para abrir el sistema de ayuda HTML con su navegador.
Escriba 'q()' para salir de R.

```

Launching java with spark-submit command /usr/local/spark-2.4.0-bin-hadoop2.7/bin/spark-submit  "sparkr-shell" /tmp/RtmpePkSOT/backend_port22961fdd5cfe
2019-01-31 14:18:20 WARN  Utils:66 - Your hostname, MBDUIV0 resolves to a loopback address: 127.0.0.1; using 156.35.119.130 instead (on interface eth2)
2019-01-31 14:18:20 WARN  Utils:66 - Set SPARK_LOCAL_IP if you need to bind to another address
2019-01-31 14:18:21 WARN  NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

```

Welcome to

```
SparkSession available as 'spark'.
```

➤

SparkR: Ejecución

- Se puede ejecutar en local o cluster
 - `sparkR.session(master="local[*]")`
 - `sparkR.session(master="yarn")`

```
> df <- as.DataFrame(faithful)
> head(df)
eruptions waiting
1      3.600      79
2      1.800      54
3      3.333      74
4      2.283      62
5      4.533      85
6      2.883      55
> head(select(df, "eruptions"))
eruptions
1      3.600
2      1.800
3      3.333
4      2.283
5      4.533
6      2.883
> quit()
Save workspace image? [y/n/c]: n
```

History Server - Mozilla Firefox

History Server x All Applications x +

inhtest1:18080 90% ... Buscar

Más visitados JobHistory

History Server

Event log directory: hdfs://inhtest1:9000/tmp/spark-logs

Last updated: 6/2/2019, 12:04:40

Show 20 entries Search:

App ID	App Name	Attempt ID	Started	Completed	Duration	Spark User	Last Updated	Event Log
local-1549450896988	SparkR		2019-02-06 11:01:35	2019-02-06 11:04:19	2.7 min	hadmin	2019-02-06 11:04:19	Download

- Crear script SparkR: (exampleSparkR.R)


```
sparkR.session(master="yarn", appName="my SparkR example")
```

```
df <- as.DataFrame(faithful)
```

```
head(df)
```

```
head(select(df, "eruptions"))
```

- Le damos permisos de ejecución: `chmod u+x exampleSparkR.R`
- Ejecución:

 hadmin@INHTEST:~/scriptsJesus

— □ ×

```
[hadmin@INHTEST scriptsJesus]$ spark-submit --master yarn --deploy-mode client exampleSparkR.R
```


SparkR: Ejecución

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
223	0	0	223	0	0 B	24 GB	0 B	3	0	0	0	0

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1523705380604_0231	hadmin	my SparkR example	SPARK	default	Tue, 12 Feb 2019 15:41:01 GMT	Tue, 12 Feb 2019 15:41:34 GMT	FINISHED	SUCCEEDED	100%	History

Se ejecutó en paralelo en dos servidores

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Logs
driver	192.168.0.150:50256	Active	0	0.0 B / 384.1 MB	0.0 B	0	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	
1	INHTEST1:50351	Active	0	0.0 B / 384.1 MB	0.0 B	1	0	0	1	1	4 s (95 ms)	0.0 B	0.0 B	0.0 B	stdout stderr
2	INHTEST2:50401	Active	0	0.0 B / 384.1 MB	0.0 B	1	0	0	2	2	5 s (0.1 s)	0.0 B	0.0 B	0.0 B	stdout stderr

- Crear un DataFrame: `read.df`
 - `path`: ubicación de los datos
 - `header`: indicamos si hay cabecera
 - `source`: tipo de datos, ej csv
 - `delimiter`: separador de columna
 - `inferSchema`: indicamos si SparkR tiene que inferir el schema
 - Se puede pasar un schema
 - `na.strings`: valor que se utiliza cuando haya NA

```
myDF <- read.df(path="/user/hadmin/Iowa_Liquor_Sales4G_cleaned.csv", source =  
"csv", delimiter=";", inferSchema = "true", header = "true", na.strings = "")
```

■ Número de columnas: colnames

```
colnames(myDF)
```

```
[1] "Invoice/Item Number" "Date"
[3] "Store Number"       "Store Name"
[5] "Address"            "City"
[7] "Zip Code"           "Store Location"
[9] "County Number"      "County"
[11] "Category"           "Category Name"
[13] "Vendor Number"      "Vendor Name"
[15] "Item Number"        "Item Description"
[17] "Pack"               "Bottle Volume (ml)"
[19] "State Bottle Cost"   "State Bottle Retail"
[21] "Bottles Sold"        "Sale (Dollars)"
[23] "Volume Sold (Liters)" "Volume Sold (Gallons)"
```

■ Primeros datos: head

□ num: número de filas a mostrar

```
head(myDF, num = 4)
```

Invoice/Item Number	Date	Store Number
1 S04016200059	2012-02-09	4139
2 S12837500006	2013-06-15	3354
3 S18405600079	2014-04-14	2588
4 S05821800010	2012-05-31	2512

	Store Name	Address
1	University Groceries	2121 UNIVERSITY AVE
2	Sam's Club 8238 / Davenport	3845 ELMORE AVE.
3	Hy-Vee Food and Drug #6 / Cedar Rapi	4035 MT. VERNON ROAD S.E.
4	Hy-Vee Wine and Spirits / Iowa City	1720 WATERFRONT DR

	City	Zip Code
1	DES MOINES	50314
2	DAVENPORT	52807
3	CEDAR RAPIDS	52403
4	IOWA CITY	52240

Store Location

1	2121 UNIVERSITY AVE DES MOINES 50314
2	3845 ELMORE AVE DAVENPORT 52807 (41.558724, -90.52708)

■ Dimensiones del dataset: `dim`, `nrow` y `ncol`

```
dim(myDF)
```

```
[1] 15542579
```

```
24
```

```
nrow(myDF)
```

```
[1] 15542579
```

```
ncol(myDF)
```

```
[1] 24
```

■ Seleccionar columnas: select

```
head(select(myDF, "Bottle Volume (ml)", "Item Number"))
```

	Bottle Volume (ml)	Item Number
1	600	65191
2	1000	23827
3	750	27410
4	1750	62400
5	1000	65257
6	750	40297

■ Filtrar: filter, where

```
myDF_filtered <- filter(myDF, myDF$'Bottle Volume (ml)' > 2000)
head(select(myDF_filtered, "Invoice/Item Number", "Bottle Volume (ml)"))
```

	Invoice/Item Number	Bottle Volume (ml)
1	S19651800003	3000
2	S11825100058	3000
3	S19322900058	3000
4	S27217800001	3000
5	S24287700005	3000
6	S27010400013	2400

- Filtrar: `filter`, `where`
 - Comprobar nulls: `isNull`

```
myDF_filtered <- filter(myDF, isNull(myDF$'Vendor Name'))  
head(select(myDF_filtered, "Invoice/Item Number", "Vendor Name"))
```

	Invoice/Item Number	Vendor Name
--	---------------------	-------------

1	INV-13688900001	
---	-----------------	--

2	INV-12438800004	
---	-----------------	--

■ Filtrar: `filter`, `where`

□ Condicionales compuestos: `&` |

```
myDF_filtered <- filter(myDF, (myDF$'Bottle Volume (ml)' > 300 & myDF$'Bottle Volume (ml)' < 500) | myDF$'Bottle Volume (ml)' > 1000 )  
head(select(myDF_filtered, "Invoice/Item Number", "Bottle Volume (ml)"))
```

	Invoice/Item Number	Bottle Volume (ml)
1	S05821800010	1750
2	S24510500021	1750
3	S09319700041	1750
4	S13147500009	1750
5	S27378700022	375
6	S15484400006	1750

- Resumen de columnas: describe, str, printSchema, schema, showDF

```
describe(myDF)
```

```
SparkDataFrame[summary:string, Invoice/Item Number:string, Store Number:string,  
Store Name:string, Address:string, City:string, Zip Code:string, Store Location:  
string, County Number:string, County:string, Category:string, Category Name:stri  
ng, Vendor Number:string, Vendor Name:string, Item Number:string, Item Descripti  
on:string, Pack:string, Bottle Volume (ml):string, State Bottle Cost:string, Sta  
te Bottle Retail:string, Bottles Sold:string, Sale (Dollars):string, Volume Sold  
(Liters):string, Volume Sold (Gallons):string]
```

Took 9 min 45 sec. Last updated by admin at February 16 2019, 7:44:03 PM.

Resumen de columnas: describe, str, printSchema, schema, showDF

```
str(myDF)

'SparkDataFrame': 24 variables:
$ Invoice/Item Number : chr "S04016200059" "S12837500006" "S18405600079" "S05821800010" "S06609500065" "S20919600062"
$ Date               : POSIXct 2012-02-09 2013-06-15 2014-04-14 2012-05-31 2012-07-17 2014-08-28
$ Store Number       : int 4139 3354 2588 2512 2555 2616
$ Store Name         : chr "University Groceries" "Sam's Club 8238 / Davenport" "Hy-Vee Food and Drug #6 / Cedar Rapi
$ Address            : chr "2121 UNIVERSITY AVE" "3845 ELMORE AVE." "4035 MT. VERNON ROAD S.E." "1720 WATERFRONT DR"
$ City              : chr "DES MOINES" "DAVENPORT" "CEDAR RAPIDS" "IOWA CITY" "KEOKUK" "CLINTON"
$ Zip Code          : chr "50314" "52807" "52403" "52240" "52632" "52732"
$ Store Location     : chr "2121 UNIVERSITY AVE DES MOINES 50314" "3845 ELMORE AVE. DAVENPORT 52807 (41.559724, -90.5
$ County Number     : int 77 82 57 52 56 23
$ County            : chr "Polk" "Scott" "Linn" "Johnson" "Lee" "Clinton"
$ Category          : int 1082900 1011100 1011200 1071100 1082900 1031200
$ Category Name     : chr "MISC. IMPORTED CORDIALS & LIQUEURS" "BLENDED WHISKIES" "STRAIGHT BOURBON WHISKIES" "AMERI
$ Vendor Number     : int 259 297 65 330 192 205
$ Vendor Name       : chr "Heaven Hill Brands" "Laird And Company" "Jim Beam Brands" "Gemini Spirits" "Sidney Frank
$ Item Number       : chr "65191" "23827" "27410" "62400" "65257" "40297"
$ Item Description  : chr "Hpnitiq Mini" "Five Star" "Jim Beam Honey" "Margaritaville Classic Lime Margarita" "Jager
$ Pack             : int 8 12 12 6 12 12
$ Bottle Volume (ml) : int 600 1000 750 1750 1000 750
$ State Bottle Cost : num 12.44 4.4 11.03 6.01 15.53 6.9
$ State Bottle Retail : num 18.66 6.59 16.55 9.52 23.3 10.35
$ Bottles Sold      : int 1 24 4 6 12 12
$ Sale (Dollars)    : num 18.66 158.16 66.2 57.12 279.6 124.2
$ Volume Sold (Liters) : num 0.6 24 3 10.5 12 9
$ Volume Sold (Gallons): num 0.16 6.34 0.79 2.77 3.17 2.38
```

- Resumen de columnas: describe, str, printSchema, schema, showDF

```
printSchema(myDF)

root
|- Invoice/Item Number: string (nullable = true)
|- Date: timestamp (nullable = true)
|- Store Number: integer (nullable = true)
|- Store Name: string (nullable = true)
|- Address: string (nullable = true)
|- City: string (nullable = true)
|- Zip Code: string (nullable = true)
|- Store Location: string (nullable = true)
|- County Number: integer (nullable = true)
|- County: string (nullable = true)
|- Category: integer (nullable = true)
|- Category Name: string (nullable = true)
|- Vendor Number: integer (nullable = true)
|- Vendor Name: string (nullable = true)
|- Item Number: string (nullable = true)
|- Item Description: string (nullable = true)
|- Pack: integer (nullable = true)
|- Bottle Volume (ml): integer (nullable = true)
|- State Bottle Cost: double (nullable = true)
|- State Bottle Retail: double (nullable = true)
|- Bottles Sold: integer (nullable = true)
|- Sale (Dollars): double (nullable = true)
|- Volume Sold (Liters): double (nullable = true)
|- Volume Sold (Gallons): double (nullable = true)
```

- Resumen de columnas: describe, str, printSchema, schema, showDF

```
showDF(describe(myDF, "Bottle Volume (ml)"))
```

```
+--+-+-----+  
|summary|Bottle Volume (ml)|  
+--+-+-----+  
|  count|          15542579|  
|   mean|    924.9787828648|  
| stddev| 695.9922246552791|  
|   min|              0|  
|   max|         378000|  
+--+-+-----+
```

■ Agregación de columnas:

□ Agrupación: `groupBy`

□ Agregación: `agg`, `sumarize` con:

- Media: `avg` y `mean`
- Máximo elemento del grupo: `max`
- Suma de los elementos del grupo: `sum`
- Primer elemento del grupo: `first`
- Último elemento del grupo: `last`
- Número de elementos: `n`
- Número de elementos distintos: `countDistinct` y `n_distinct`
- Desviación estándar: `sd`, `stddev`, `stddev_samp` y `stddev_pop`
- Varianza: `var`, `variance`, `var_samp` y `var_pop`
- Forma de la distribución: `kurtosis`, `skewness`

■ Agregación de columnas:

```
aggregationCountySales <- agg(groupBy(myDF, myDF$'County'), "min of sales (dollars)" = min(myDF$'Sale (Dollars)'), "max of sales (dollars)" = max(myDF$'Sale (Dollars)'))
collect(aggregationCountySales)
```

	County	min of sales (dollars)	max of sales (dollars)
--	--------	------------------------	------------------------

1	HANCOCK	0.00	3753.00
2	JOHNSON	0.00	69498.00
3	ADAIR	0.00	2239.50
4	EMMET	0.00	16200.00
5	Harrison	0.00	3354.12
6	Scott	0.00	34619.76
7	Lucas	0.00	12150.00
8	Monroe	0.00	9630.00

- Último elemento del grupo: last
- Forma de la distribución: kurtosis, skewness

- Resumen estadístico: `approxQuantile`
 - Nos da una aproximación que puede tener un error
 - `relativeError`

```
approxQuantile(myDF, col = "Sale (Dollars)", probabilities = c(0, 0.25, 0.5, 0.75, 1), relativeError = 0.01)

[[1]]
[1] 0

[[2]]
[1] 31.48

[[3]]
[1] 70.56

[[4]]
[1] 141.72

[[5]]
[1] 279557.3
```


- Resumen estadístico varias columnas:
 - Covarianzas: `cov`, `covar_samp` y `covar_pop`
 - Correlaciones: `corr`
 - Tablas cruzadas: `crosstab`

```
cov(myDF, "Bottles Sold", "Sale (Dollars)")
```

```
[1] 10095.62
```

- Resumen estadístico varias columnas:
 - Covarianzas: `cov`, `covar_samp` y `covar_pop`
 - Correlaciones: `corr`
 - Tablas cruzadas: `crosstab`

```
corr(myDF, "Bottles Sold", "Sale (Dollars)")  
  
[1] 0.8449618
```

- Resumen estadístico varias columnas:
 - Covarianzas: `cov`, `covar_samp` y `covar_pop`
 - Correlaciones: `corr`
 - Tablas cruzadas: `crosstab`

```
crosstab(myDF, "City", "Category")
```

	City_Category	1011100	1011200	1011250	1011300	1011400	1011500
1	NEW SHARON	19	34	0	22	0	3
2	FREDERICKSBURG	184	57	0	33	0	10
3	WALKER	14	24	0	24	0	15
4	LAKE CITY	291	135	0	54	0	21
5	ZWITNGLE	130	27	0	65	0	0

■ Ordenar columna: arrange, orderBy

□ Le indicamos asc o desc

```
myDF_sorted <- arrange(myDF, desc(myDF$'Sale (Dollars)'))  
head(select(myDF_sorted, "Invoice/Item Number", "Sale (Dollars)"))
```

	Invoice/Item Number	Sale (Dollars)
1	INV-14774700005	279557.3
2	S09275100052	254100.0
3	S05867400001	254100.0
4	S12933100005	196004.9
5	S31923800002	181962.0
6	S09484600001	116094.0

■ Añadir columna: withColumn

```
myDF_extraColumn <- withColumn(myDF, "Bottle Volume (l)", myDF$'Bottle Volume (ml)' / 1000)  
head(select(myDF_extraColumn, "Invoice/Item Number", "Bottle Volume (ml)", "Bottle Volume (l)"))
```

	Invoice/Item Number	Bottle Volume (ml)	Bottle Volume (l)
1	S04016200059	600	0.60
2	S12837500006	1000	1.00
3	S18405600079	750	0.75
4	S05821800010	1750	1.75
5	S06609500065	1000	1.00
6	S20919600062	750	0.75

■ Eliminar columna: NULL

```
myDFCopy <- myDF
myDFCopy$'Bottle Volume (ml)' <- NULL
myDFCopy$'Invoice/Item Number' <- NULL
myDFCopy$'Store Number' <- NULL
myDFCopy$'Address' <- NULL
myDFCopy$'Zip Code' <- NULL
myDFCopy$'County Number' <- NULL
myDFCopy$'Category' <- NULL
myDFCopy$'Vendor Number' <- NULL
myDFCopy$'Item Number' <- NULL
myDFCopy$'Pack' <- NULL
myDFCopy$'State Bottle Retail' <- NULL
myDFCopy$'Sale (Dollars)' <- NULL
myDFCopy$'Volume Sold (Gallons)' <- NULL
myDFCopy$'Date' <- NULL
myDFCopy$'Store Name' <- NULL
```

```
cat("Num cols myDF with all columns:", ncol(myDF), "\n")
cat("Num cols myDF without all columns:", ncol(myDFCopy), "\n")
cat("Columns:\n")
columns(myDFCopy)
```

Num cols myDF with all columns: 24

Num cols myDF without all columns: 9

Columns:

[1] "City"	"Store Location"	"County"
[4] "Category Name"	"Vendor Name"	"Item Description"
[7] "State Bottle Cost"	"Bottles Sold"	"Volume Sold (Liters)"

■ Añadir filas:

□ Unión de dos DataFrames

□ rbind

```
newRow <- data.frame("Invoice/Item Number" = "S93034200059",  
  "Date" = as.POSIXct("2018-08-30"),  
  "Store Number" = as.integer(3354),  
  "Store Name" = "Sam's Club 8238 / Davenport",  
  "Address" = "3845 ELMORE AVE.",  
  "City" = "DAVENPORT",  
  "Zip Code" = "52807",  
  "Store Location" = "3845 ELMORE AVE. DAVENPORT 52807 (41.559724, -90.52708)",  
  "County Number" = as.integer(82),  
  "County" = "Scott",  
  "Category" = as.integer(1031080),  
  "Category Name" = "VODKA 80 PROOF",  
  "Vendor Number" = as.integer(297),  
  "Vendor Name" = "Laird And Company",  
  "Item Number" = "35917",  
  "Item Description" = "Five O'clock Vodka",  
  "Pack" = as.integer(12),  
  "Bottle Volume (ml)" = as.integer(1000),  
  "State Bottle Cost" = 4.39,  
  "State Bottle Retail" = 6.83,  
  "Bottles Sold" = as.integer(12),  
  "Sale (Dollars)" = 300.48,  
  "Volume Sold (Liters)" = 48,  
  "Volume Sold (Gallons)" = 12.68)  
  
myRowDF <- as.DataFrame(newRow)  
  
myDFCopy <- rbind(myDF, myRowDF)
```

```
nrow(myDF)  
nrow(myDFCopy)
```

```
[1] 15542579  
[1] 15542580
```

- Obtener datos únicos (conjunto de datos):
unique

```
myDFCountyUnique <- unique(select(myDF, "County"))  
head(myDFCountyUnique)
```

County

```
1 HANCOCK  
2 JOHNSON  
3 ADAIR  
4 EMMET  
5 Harrison  
6 Scott
```


■ Muestreo:

□ Big Data: `sample`

□ Datos pequeños: `takeSample`

```
myDFSampled <- sample(myDF, withReplacement = FALSE, fraction = 0.25, seed = 1)
nrow(myDF)
nrow(myDFSampled)
```

```
[1] 15542579
```

```
[1] 3884891
```

- Eliminar valores nulos: dropNa
 - cols: columnas en las que analizar Nulls
 - how: condiciones para eliminar la fila:
 - all: se elimina la fila si todas sus cols tienen Nulls
 - any: se elimina la fila si alguna de sus cols tiene NULL
 - minNonNulls: Número mínimo de valores no nulos posibles en la fila

```
myDFWithLessNa <- dropna(myDF, cols = list("Bottles Sold", "Vendor Name"))  
nrow(myDF)  
nrow(myDFWithLessNa)
```

```
[1] 15542579
```

```
[1] 15542577
```

- Eliminar valores nulos: dropNa
 - cols: columnas en las que analizar Nulls
 - how: condiciones para eliminar la fila:
 - all: se elimina la fila si todas sus cols tienen Nulls
 - any: se elimina la fila si alguna de sus cols tiene NULL
 - minNonNulls: Número mínimo de valores no nulos posibles en la fila

```
myDFWithLessNa <- dropna(myDF, cols = list("Bottles Sold", "Vendor Name"), how = "all")  
nrow(myDFWithLessNa)
```

```
[1] 15542579
```

- Eliminar valores nulos: dropNa
 - cols: columnas en las que analizar Nulls
 - how: condiciones para eliminar la fila:
 - all: se elimina la fila si todas sus cols tienen Nulls
 - any: se elimina la fila si alguna de sus cols tiene NULL
 - minNonNulls: Número mínimo de valores no nulos posibles en la fila

```
myDFWithLessNa <- dropna(myDF, cols = list("Bottles Sold", "Vendor Name", "Store Location"), minNonNulls = 2)  
nrow(myDFWithLessNa)
```

```
[1] 15542579
```

- Filtrar/cambiar valores nulos: `fillNA`
 - `cols`: columnas a analizar
 - `value`: valor por el que se sustituirá el NULL

```
myDFWithLessNa <- fillna(myDF, cols = list("Store Location", "Vendor Name"), value = "Unknown")  
nrow(filter(myDF, myDF$'Vendor Name' == "Unknown"))  
nrow(filter(myDFWithLessNa, myDFWithLessNa$'Vendor Name' == "Unknown"))
```

```
[1] 0
```

```
[1] 2
```

■ dapply

- Aplica una función a cada partición
 - Recibe la partición como un data.frame de R
- ~ mapPartitions
- Genera un data.frame de R que se transforma en la partición
- Operación costosa: serializa y deserializa información entre JVM y R

■ gapply

□ Función ejecutada por cada agrupación

- Columna para agrupar
- Función para ejecutar en cada grupo

```
mySchema <- structType(structField("Bottles Sold", "integer"), structField("Max Sale (Dollars)", "double"))
maxSalesByNumOfBottlesSold <- gapply(
  myDF,
  "Bottles Sold",
  function(key, x) {
    y <- data.frame(key, max(x$'Sale (Dollars)'))
  },
  mySchema)
head(maxSalesByNumOfBottlesSold)
```

	Bottles Sold	Max Sale (Dollars)
1	148	1593.96
2	243	6011.82
3	540	111277.80
4	1522	4870.40
5	31	1468.47
6	516	14226.12

■ lapply

- Ejecutar varias tareas en paralelo
- Tiene que ser capaz de ejecutar cada tarea en un único servidor
 - Ejemplo: entrenar varios modelos con diferentes parámetros
 - Crear lista con los parámetros del modelo
 - lapply entrena en un servidor un modelo, en otro servidor otro modelo, ...
 - Tiene que ser capaz de entrenar cada modelo en un único servidor

- Unión de columnas: `join`
 - DataFrames a unir
 - `joinExpr`: Columna que queremos unir
 - `joinType`: `inner`, `outer`, `full`, `fullouter`, `leftouter`, `left_outer`, `left`, `right_outer`, `rightouter`, `right`, y `leftsemi`

```
df_r1 <- data.frame("Client" = c("Alice", "Bob", "Carol", "Dave"),
                    "Community" = c("AST", "AST", "GAL", "CAN"))
DF1 <- as.DataFrame(data = df_r1)

df_r2 <- data.frame("Community" = c("AST", "CAN", "GAL", "CAST-L"),
                    "Population" = c(1028244, 580229, 2708339, 1028244))
DF2 <- as.DataFrame(data = df_r2)

joinDF1_2 <- join(DF1, DF2, DF1$Community == DF2$Community, joinType = "inner")
```

[1] "DF1:"

	Client	Community
1	Alice	AST
2	Bob	AST
3	Carol	GAL
4	Dave	CAN

[1] "DF2"

	Community	Population
1	AST	1028244
2	CAN	580229
3	GAL	2708339
4	CAST-L	1028244

[1] "inner join:"

	Client	Community	Community	Population
1	Alice	AST	AST	1028244
2	Bob	AST	AST	1028244
3	Carol	GAL	GAL	2708339
4	Dave	CAN	CAN	580229

■ Unión de filas: rbind

□ Añade filas de un DataFrame a otro

```
df_r1 <- data.frame("Client" = c("Alice", "Bob", "Carol", "Dave"),  
                    "Community" = c("AST", "AST", "GAL", "CAN"))  
DF1 <- as.DataFrame(data = df_r1)  
  
df_r2 <- data.frame("Client" = c("Eva", "Frank"),  
                    "Community" = c("CAN", "AST"))  
DF2 <- as.DataFrame(data = df_r2)  
  
unionDF1_2 <- rbind(DF1, DF2)
```

```
[1] "DF1:"  
Client Community  
1 Alice      AST  
2 Bob       AST  
3 Carol     GAL  
4 Dave      CAN
```

```
[1] "DF2:"  
Client Community  
1 Eva       CAN  
2 Frank     AST
```

```
[1] "Union of rows:"  
Client Community  
1 Alice      AST  
2 Bob       AST  
3 Carol     GAL  
4 Dave      CAN  
5 Eva       CAN  
6 Frank     AST
```

■ Intersección: intersect

□ Obtenemos las filas que están en ambos DataFrames

```
df_r1 <- data.frame("Client" = c("Alice", "Bob", "Carol", "Dave"),  
                    "Community" = c("AST", "AST", "GAL", "CAN"))  
DF1 <- as.DataFrame(data = df_r1)  
  
df_r2 <- data.frame("Client" = c("Eva", "Frank", "Alice", "Dave", "Carol"),  
                    "Community" = c("CAN", "AST", "AST", "AST", "GAL"))  
DF2 <- as.DataFrame(data = df_r2)  
  
intersectionDF1_2 <- intersect(DF1, DF2)
```

[1] "DF1:"

	Client	Community
1	Alice	AST
2	Bob	AST
3	Carol	GAL
4	Dave	CAN

[1] "DF2:"

	Client	Community
1	Eva	CAN
2	Frank	AST
3	Alice	AST
4	Dave	AST
5	Carol	GAL

[1] "Intersection:"

	Client	Community
1	Alice	AST
2	Carol	GAL

- Resta de filas: `except`
 - Quitar las filas que estén en otro DataFrame
 - Puede utilizarse para dividir el dataset en train y test
 - Train: con un muestreo
 - Test: resta del DataFrame original menos el train
 - Para dividir un DataFrame en train y test es mejor utilizar `randomSplit` (a partir de la versión 2.0 de Spark)

```
df_r1 <- data.frame("Client" = c("Alice", "Bob", "Carol", "Dave"),  
                    "Community" = c("AST", "AST", "GAL", "CAN"))  
DF1 <- as.DataFrame(data = df_r1)  
  
df_r2 <- data.frame("Client" = c("Eva", "Frank", "Alice", "Dave", "Carol"),  
                    "Community" = c("CAN", "AST", "AST", "AST", "GAL"))  
DF2 <- as.DataFrame(data = df_r2)  
  
DF1_substract_DF2 <- except(DF1, DF2)
```

```
[1] "DF1:"  
Client Community  
1 Alice      AST  
2 Bob        AST  
3 Carol      GAL  
4 Dave       CAN
```

```
[1] "DF2:"  
Client Community  
1 Eva        CAN  
2 Frank      AST  
3 Alice      AST  
4 Dave       AST  
5 Carol      GAL
```

```
[1] "DF1 subtraction DF2:"  
Client Community  
1 Dave       CAN  
2 Bob        AST
```

■ SQL

- Crear tabla temporal: `registerTempTable`
- Consulta: `sql`

```
registerTempTable(myDF, "myTable")  
myDF_afterQuery <- sql("SELECT * FROM myTable WHERE `Pack` < 10")  
head(select(myDF_afterQuery, "Invoice/Item Number", "Pack"))
```

	Invoice/Item Number	Pack
1	S04016200059	8
2	S05821800010	6
3	S24510500021	6
4	S09319700041	6
5	S13147500009	6
6	S08233300074	6

■ Sustituir datos: ifelse

- Crear una columna con un valor u otro dependiendo de una condición
 - Condición que se evalúa por cada dato
 - Yes: valor por el que se debe sustituir en caso de que el dato sea sí
 - No: idem para el no

```
myDF$"Pack size" <- ifelse(test = myDF$"Pack" < 10, yes = "Few", no = myDF$"Pack")  
head(select(myDF, "Invoice/Item Number", "Pack", "Pack size"))
```

	Invoice/Item Number	Pack	Pack size
1	S04016200059	8	Few
2	S12837500006	12	12
3	S18405600079	12	12
4	S05821800010	6	Few
5	S06609500065	12	12
6	S20919600062	12	12

- Guardar DataFrame: `write.df`
 - DataFrame
 - path
 - source: tipo de dato para que lo guarde en ese formato
 - mode: tipo de guardado:
 - overwrite: si hay un archivo lo elimina
 - append: si hay un archivo, guarda el DataFrame a continuación
 - error: si hay un archivo, salta una excepción
 - ignore: si hay un archivo, no guarda nada.
 - Sólo se guarda el DataFrame si no existe nada

■ Cacheo y descacheo: cache y unpersist

```
cache(myDF)
```

```
SparkDataFrame[Invoice/Item Number:string, Date:timestamp, Store Number:int, Store Name:string, Address:string, City:string, Zip Code:string, Store Location:string, County Number:int, County:string, Category:int, Category Name:string, Vendor Number:int, Vendor Name:string, Item Number:string, Item Description:string, Pack:int, Bottle Volume (ml):int, State Bottle Cost:double, State Bottle Retail:double, Bottles Sold:int, Sale (Dollars):double, Volume Sold (Liters):double, Volume Sold (Gallons):double, Pack size:string]
```

```
unpersist(myDF)
```

```
SparkDataFrame[Invoice/Item Number:string, Date:timestamp, Store Number:int, Store Name:string, Address:string, City:string, Zip Code:string, Store Location:string, County Number:int, County:string, Category:int, Category Name:string, Vendor Number:int, Vendor Name:string, Item Number:string, Item Description:string, Pack:int, Bottle Volume (ml):int, State Bottle Cost:double, State Bottle Retail:double, Bottles Sold:int, Sale (Dollars):double, Volume Sold (Liters):double, Volume Sold (Gallons):double, Pack size:string]
```


- Crear data.frame: collect
 - Sólo cuando tenemos pocos datos
 - data.frame de R no soporta grandes cantidades de datos

```
%spark.r
myDFSmall <- describe(myDF, "Bottle Volume (ml)")
myDFSmall_r <- collect(myDFSmall)

print("DataFrame Spark:")
typeof(myDFSmall)

print("dataframe R:")
typeof(myDFSmall_r)
myDFSmall_r
```

```
[1] "DataFrame Spark:"
[1] "S4"
```

```
[1] "dataframe R:"
[1] "list"
summary Bottle Volume (ml)
1  count      15542579
2  mean       924.9787828648
3  stddev   695.9922246552791
4   min                0
5   max       378000
```

- Visualización de datos

- Es difícil con grandes cantidades de Datos

- SparkR: librería `ggplot2.SparkR`

- Utiliza DataFrames de Spark

- Sólo disponible en versiones antiguas de Spark

- Extrae todo lo que necesita el gráfico utilizando Spark y luego genera el gráfico con `ggplot` de R

■ Boxplot

- Se necesitan 5 datos: mínimo, primer cuartil, mediana, tercer cuartil y máximo
- Pasos:
 - Transformar el DataFrame a un DataFrame con los 5 datos
 - Convertirlo en `data.frame` de R
 - Crear el gráfico con las librerías de R

■ Boxplot

```
mySchemaSummary <- structType(structField("min", "double"),
                                structField("Quantile25", "double"),
                                structField("median", "double"),
                                structField("Quantile75", "double"),
                                structField("max", "double"))

myDFSSummary <- as.DataFrame(data = data.frame('min' = double(), "Quantile25" = double(), "median"
= double(), "Quantile75" = double(), "max" = double()), schema = mySchemaSummary)

myDFQuantiles <- approxQuantile(myDF, col = "Sale (Dollars)", probabilities = c(0, 0.25, 0.5, 0.75
, 1), relativeError = 0.01)


newRow <- data.frame("min" = myDFQuantiles[[1]],
                    "Quantile25" = myDFQuantiles[[2]],
                    "median" = myDFQuantiles[[3]],
                    "Quantile75" = myDFQuantiles[[4]],
                    "max" = myDFQuantiles[[5]])

myRowDF <- as.DataFrame(newRow)

myDFSSummary <- rbind(myDFSSummary, myRowDF)

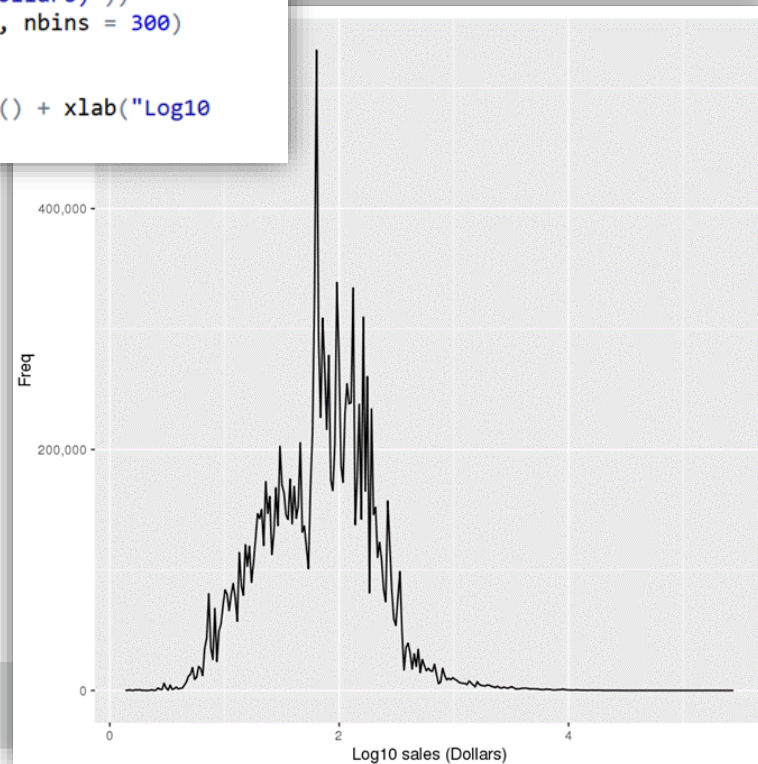
myDFQuantiles_r <- collect(myDFSSummary)

library(ggplot2)
library(scales)
ggplot(myDFQuantiles_r, aes(1)) + geom_boxplot(aes(ymin = min, lower = Quantile25, middle = median
, upper = Quantile75, ymax = max), stat = "identity") + scale_y_continuous(trans='log2',
labels = dollar, name = "log2 sales (dollar)") + scale_x_continuous(labels = c(), name = "")
```



■ Histograma

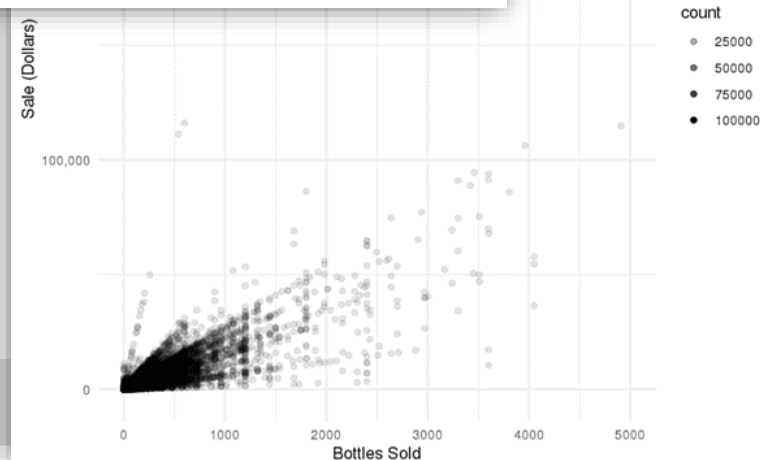
```
myDFWithLogSales <- withColumn(myDF, "Log sale (Dollars)", log10(myDF$'Sale (Dollars)'))  
myDFSummarized_r_histSales <- histogram(myDFWithLogSales, "Log sale (Dollars)", nbins = 300)  
library(ggplot2)  
library(scales)  
ggplot(myDFSummarized_r_histSales, aes(x = centroids, y = counts)) + geom_path() + xlab("Log10  
sales (Dollars)") + scale_y_continuous(labels = comma, name = "Freq")
```



■ Dispersión

```
myDFBottlesSoldSale <- select(myDF, "Bottles Sold", "Sale (Dollars)")
myDFBottlesSoldSaleCount <- agg(groupBy(myDF, myDF$'Bottles Sold', myDF$'Sale (Dollars)'), count =
  n(myDF$'Bottles Sold'))

myDFBottlesSoldSaleCount_r <- collect(myDFBottlesSoldSaleCount)
library(ggplot2)
library(scales)
ggplot(myDFBottlesSoldSaleCount_r, aes(x = `Bottles Sold`, y = `Sale (Dollars)`, alpha = `count`))
  + geom_point() + scale_x_continuous(lim = c(0,5000), name = "Bottles Sold") +
  scale_y_continuous(labels = comma, name = "Sale (Dollars)") + theme_minimal()
```



- No todos los modelos de ML son paralelizables
- Spark:
 - Spark ML y MLlib
 - Mahout: <https://mahout.apache.org/>
 - ...

■ SparkR

□ Classification:

- `spark.logit`: Logistic Regression
- `spark.mlp`: Multilayer Perceptron (MLP)
- `spark.naiveBayes`: Naive Bayes
- `spark.svmLinear`: Linear Support Vector Machine
- `spark.fmClassifier`: Factorization Machines classifier

□ Regression

- `spark.survreg`: Accelerated Failure Time (AFT) Survival Model
- `spark.glm` o `glm`: Generalized Linear Model (GLM)
- `spark.isoreg`: Isotonic Regression
- `spark.lm`: Linear Regression
- `spark.fmRegressor`: Factorization Machines regressor

Obtenido de la página oficial de Spark (Actualizado a 12/01/2022):
<https://spark.apache.org/docs/latest/sparkr.html#machine-learning>

□ Tree

- `spark.decisionTree`: Decision Tree for Regression and Classification
- `spark.gbt`: Gradient Boosted Trees for Regression and Classification
- `spark.randomForest`: Random Forest for Regression and Classification

□ Clustering

- `spark.bisectingKmeans`: Bisecting k-means
- `spark.gaussianMixture`: Gaussian Mixture Model (GMM)
- `spark.kmeans`: K-Means
- `spark.lda`: Latent Dirichlet Allocation (LDA)
- `spark.powerIterationClustering` (PIC): Power Iteration Clustering (PIC)

□ Collaborative Filtering

- `spark.als`: Alternating Least Squares (ALS)

□ Frequent Pattern Mining

- `spark.fpGrowth`: FP-growth
- `spark.prefixSpan` : PrefixSpan

□ Statistics

- `spark.kstest`: Kolmogorov-Smirnov Test

■ Modelo de regresión lineal

```
myDF_list <- randomSplit(dropna(select(myDF, "Bottles Sold", "Sale (Dollars)")), c(7,3), 2)
trainDF <- myDF_list[[1]]
testDF <- myDF_list[[2]]

colnames(trainDF) <- c("Bottles_Sold", "Sale_Dollars")
colnames(testDF) <- c("Bottles_Sold", "Sale_Dollars")

myModel <- glm(Sale_Dollars ~ Bottles_Sold, data = trainDF, family = gaussian)
print("my model:")
summary(myModel)
```

[1] "my model:"

Deviance Residuals:

(Note: These are approximate quantiles with relative error <= 0.01)

Min	1Q	Median	3Q	Max
-54695	-16	6	29	107556

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-11.872	0.078176	-151.86	0
Bottles_Sold	14.249	0.0027604	5162	0

(Dispersion parameter for gaussian family taken to be 58125.52)

Null deviance: 2.1812e+12 on 10878563 degrees of freedom

Residual deviance: 6.3232e+11 on 10878562 degrees of freedom
AIC: 150213795

Number of Fisher Scoring iterations: 1

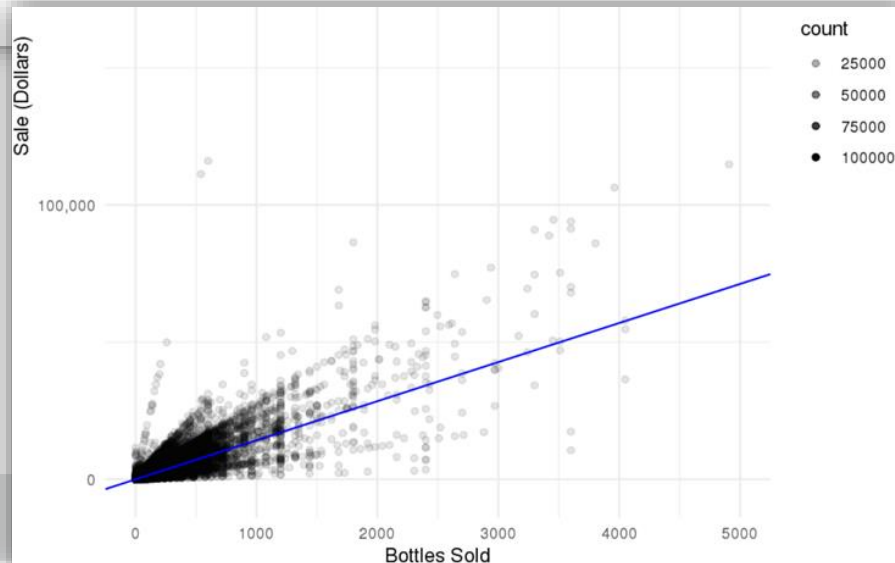
■ Modelo de regresión lineal

```
myPredictions <- predict(myModel, testDF)
print("predictions")
head(myPredictions)
```

```
[1] "predictions"
  Bottles_Sold Sale_Dollars label prediction
1           1           0     0   2.376989
2           1           0     0   2.376989
3           1           0     0   2.376989
4           1           0     0   2.376989
5           1           0     0   2.376989
6           1           0     0   2.376989
```

■ Modelo de regresión lineal

```
myModelIntercept <- as.numeric(summary(myModel)$coefficients["(Intercept)",1])
myModelSlope <- as.numeric(summary(myModel)$coefficients["Bottles_Sold",1])
ggplot(myDFBottlesSoldSaleCount_r, aes(x = `Bottles Sold`, y = `Sale (Dollars)`, alpha
  = `count`)) + geom_point() + scale_x_continuous(lim = c(0,5000), name = "Bottles
  Sold") + scale_y_continuous(labels = comma, name = "Sale (Dollars)") +
  theme_minimal() + geom_abline(intercept = myModelIntercept, slope = myModelSlope,
  colour = "blue")
```



■ Modelo de regresión lineal

- Rsquared: (código obtenido de https://github.com/UrbanInstitute/sparkr-tutorials/blob/master/09_glm.md)

```
myTrainPredictions <- predict(myModel, trainDF)

y <- myTrainPredictions$Sale_Dollars
y_avg <- collect(agg(myTrainPredictions, y_avg = mean(y)))$y_avg
myTrainPredictionsTransformed <- transform(myTrainPredictions, y_hat =
  myTrainPredictions$prediction, sq_res = (y - myTrainPredictions$prediction)^2,
  sq_tot = (y - y_avg)^2, res = y - myTrainPredictions$prediction)
myTrainPredictionsTransformed$prediction <- NULL
head(myTrainPredictionsTransformed)

SSR <- collect(agg(myTrainPredictionsTransformed, SSR = sum
  (myTrainPredictionsTransformed$sq_res)))
SST <- collect(agg(myTrainPredictionsTransformed, SST = sum
  (myTrainPredictionsTransformed$sq_tot)))

Rsqr <- 1-(SSR[[1]]/SST[[1]])
p <- 10
N <- nrow(myTrainPredictionsTransformed)
aRsqr <- Rsqr - (1 - Rsqr)*((p - 1)/(N - p))
```

```
[1] "Rsqr:"
[1] 0.7100977
[1] "aRsqr:"
[1] 0.7100974
```

- R vs SparkR
- Instalación SparkR
- Operaciones SparkR
- Visualización de datos
- Machine Learning

Gracias